

Brain Tumour Detection Using Convolutional Neural Networks

CEP: Artificial Intelligence and Machine Learning Lab

Muhammad Ahmad (FA23-BCE-113)

Uzair (FA23-BCE-098)

Computer Engineering Department
COMSATS University Islamabad – Lahore Campus

May 19, 2026

Outline

- 1 Introduction
- 2 Dataset
- 3 CNN Architecture
- 4 Training Procedure
- 5 Results
- 6 Mathematical Foundation
- 7 Sustainable Development Goals
- 8 Future Enhancements
- 9 Conclusions

Problem Statement

- **Challenge:** Brain tumours are among the most life-threatening neurological disorders
- Manual MRI interpretation is time-consuming and requires specialist radiologists
- Human analysis is prone to fatigue-induced errors and inter-observer variability
- Limited radiologist access in resource-constrained and rural regions
- **Solution:** Deep CNN-based automated binary classification of brain MRI scans

*Objective: Train a robust CNN to classify brain MRI scans as **Normal** or **Tumour-Positive** with clinical-grade sensitivity*

Key Objectives

Primary Goals:

- Binary classification of MRI scans
- Maximise sensitivity (minimise false negatives)
- Learn discriminative spatial features
- Enable real-time inference via API
- Support downstream DIP pipeline

Accuracy

Sensitivity

Speed

Scalability

Technical Requirements:

- Four-stage CNN architecture
- Data augmentation for generalisation
- Quantitative metric evaluation
- PCA feature space visualisation
- FastAPI deployment-ready export

Dataset Overview


Brain Tumour MRI Dataset:

- 255 grayscale MRI scan images
- Two classes: Normal and Tumour
- Resized to 128×128 pixels
- Label mode: binary
- 80/20 train-validation split


Loading Method:

```
tf.keras.utils.image_dataset_from  
_directory() with validation_split=0.20
```

Class Distribution




Data Augmentation Strategy



ImageDataGenerator — applied only to training split

Four-Stage CNN Architecture




Layer-by-Layer Summary

black			
Input ($128 \times 128 \times 3$)	(None,128,128,3)	—	—
Conv2D (16, 3×3)	(None,126,126,16)	448	ReLU
MaxPool (2×2)	(None, 63, 63, 16)	0	—
Conv2D (32, 3×3)	(None, 61, 61, 32)	4,640	ReLU
MaxPool (2×2)	(None, 30, 30, 32)	0	—
Conv2D (64, 3×3)	(None, 28, 28, 64)	18,496	ReLU
MaxPool (2×2)	(None, 14, 14, 64)	0	—
Conv2D (128, 3×3)	(None, 12, 12,128)	73,856	ReLU
MaxPool (2×2)	(None, 6, 6,128)	0	—
Flatten	(None, 4,608)	0	—
Dense (256)	(None, 256)	1,179,904	ReLU
Dense (1) — Output	(None, 1)	257	Sigmoid
gray!15 Total Parameters	1,277,601 (4.87 MB)		

Training Configuration

Hyperparameters:



- **Optimizer:** Adam ($lr = 0.001$)
- **Loss:** BinaryCrossentropy
- **Metric:** Accuracy
- **Epochs:** 10 (EarlyStopping, patience = 5)
- **Batch Size:** 32
- **Val Split:** 20 % (51 images)



Epoch-by-Epoch Training Metrics

black					
1	0.7071	54.51 %	0.6383	60.78 %	
2	0.5991	65.88 %	0.4846	82.35 %	
3	0.5103	74.90 %	0.6240	60.78 %	
4	0.5112	74.51 %	0.4571	82.35 %	
5	0.5134	76.47 %	0.4079	80.39 %	
6	0.4735	78.82 %	0.3475	84.31 %	
7	0.4349	81.96 %	0.3224	92.16 %	
8	0.3916	83.14 %	0.2528	88.24 %	
9	0.3483	83.53 %	0.2160	94.12 %	
green!12 10	0.2941	86.27 %	0.1850	94.12 %	Best

Training and Validation Curves



Model Performance — Key Metrics

Test Acc**92.16%**


Val Acc**94.12%**

Sensitivity**96.18%**

F1-Score**93.79%**

black		
True Positives (Tumour)	151	Correct tumour detections
True Negatives (Normal)	84	Correct normal classifications
False Positives	14	Normal misclassified as Tumour
False Negatives	6	Tumour misclassified as Normal
Precision	91.52 %	$TP / (TP + FP)$
Specificity	85.71 %	$TN / (TN + FP)$

Confusion Matrix



Confusion Matrix — Clinical Interpretation

What the numbers mean:

- **151 TP:** Tumours correctly flagged for treatment
- **84 TN:** Healthy patients correctly cleared
- **14 FP:** Unnecessary follow-up scans (low risk)
- **6 FN:** Missed tumours — most critical error


TP = 151	FN=6
----------	------

← minimise this

Clinical Context:

- Sensitivity 96.18 % — 96 of 100 tumours detected
- Suitable as a preliminary screening tool
- False positives trigger radiologist review (safe)
- False negatives bypass treatment (dangerous)
- Model is intentionally tuned to minimise FN

PCA Feature Space Visualisation



What PCA reveals:

- Features extracted from penultimate Dense(256) layer
- 256-dimensional vectors projected to 2D
- PC1 explains **75.3 %** of variance
- PC2 explains **15.7 %** of variance
- Total: **90.9 %** variance retained

Significance:

- Clear spatial separation between classes
- Confirms CNN learned discriminative features
- Not relying on superficial pixel patterns
- High separability → reliable decisions



Loss Function and Optimisation

Binary Cross-Entropy Loss:

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N \left[y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i) \right]$$

Adam Optimiser Update Rule:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t \quad \text{where } \eta = 0.001$$

Adaptive learning rate

Momentum correction

Per-parameter update


Precision, Recall, F1-Score:

$$\text{Precision} = \frac{TP}{TP + FP} = 91.52\% \quad \text{Recall} = \frac{TP}{TP + FN} = 96.18\%$$


$$F_1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} = 93.79\%$$

Sigmoid Output (Binary Decision):

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}} \Rightarrow \text{class} = \begin{cases} \text{Tumour} & \hat{y} \geq 0.5 \\ \text{Normal} & \hat{y} < 0.5 \end{cases}$$



Future Roadmap



Key Achievements

Technical Success:

- 4-stage CNN trained from scratch
- 92.16 % test accuracy achieved
- 94.12 % validation accuracy (epoch 10)
- 96.18 % sensitivity — only 6 FN in 157 tumours
- PCA confirms 90.9 % class separability
- Exported .keras model via FastAPI

Clinical and Social Impact:


- Pre-screening tool for radiologists
- Upstream classifier for DIP pipeline
- Ollama LLM chatbot for patient Q&A
- Deploys without GPU or cloud
- Advances SDG 3, 9, and 10



Questions and Discussion

Repository: `github.com/Muhammad-Ahmad17/cancer_cells`

Contact: `muhammadahmad171105@gmail.com`



Supervised by: **Engr. Abubakar Talha**